

Find a duplicate, *Space Edition*[™].

We have a vector of integers, where:

1. The integers are in the range $1..n$
2. The vector has a length of $n + 1$

It follows that our vector has *at least* one integer which appears *at least* twice. But it may have *several* duplicates, and each duplicate may appear *more than* twice.

Write a function which finds an integer that appears more than once in our vector. (If there are multiple duplicates, you only need to find one of them.)

We're going to run this function on our new, super-hip MacBook Pro With Retina Display[™]. Thing is, the damn thing came with the RAM soldered right to the motherboard, so we can't upgrade our RAM. **So we need to optimize for space!**

Gotchas

We can do this in $O(1)$ space.

We can do this in less than $O(n^2)$ time while keeping $O(1)$ space.

We can do this in $O(n \lg n)$ time and $O(1)$ space.

We can do this *without* destroying the input.

Most $O(n \lg n)$ algorithms double something or cut something in half. How can we rule out half of the numbers each time we iterate through the vector?

Breakdown

This one's a classic! We just do one walk through the vector, using a set to keep track of which items we've seen!

```
unsigned int findRepeat(const vector<unsigned int>& numbers)
{
    unordered_set<unsigned int> numbersSeen;
    for (unsigned int number : numbers) {
        auto it = numbersSeen.find(number);
        if (it != numbersSeen.end()) {
            return number;
        }
        else {
            numbersSeen.insert(number);
        }
    }

    // whoops--no duplicate
    throw invalid_argument("no duplicate!");
}
```

C++ ▼

Bam. $O(n)$ time and ... $O(n)$ space ...

Right, we're supposed to optimize for *space*. $O(n)$ is actually kinda high space-wise. Hm. We can probably get $O(1)$...

We can "brute force" this by taking each number in the range $1..n$ and, for each, walking through the vector to see if it appears twice.

```

unsigned int findRepeat(const vector <unsigned int>& numbers)
{
    for (unsigned int needle = 1; needle < numbers.size(); ++needle) {
        bool hasBeenSeen = false;
        for (unsigned int number : numbers) {
            if (number == needle) {
                if (hasBeenSeen) {
                    return number;
                }
                else {
                    hasBeenSeen = true;
                }
            }
        }
    }

    // whoops--no duplicate
    throw invalid_argument("no duplicate!");
}

```

This is $O(1)$ space and $O(n^2)$ time.

That space complexity can't be beat, but the time cost seems a bit high. Can we do better?

One way to beat $O(n^2)$ time is to get $O(n \lg n)$ time. Sorting takes $O(n \lg n)$ time. And if we sorted the vector, any duplicates would be right next to each-other!

But if we start off by sorting our vector we'll need to take $O(n)$ space to store the sorted vector...

...unless we sort the input vector in place!

Okay, so this'll work:

1. Do an in-place sort of the vector (for example an in-place merge sort).
2. Walk through the now-sorted vector from left to right.
3. Return as soon as we find two adjacent numbers which are the same.

This'll keep us at $O(1)$ space and bring us down to $O(n \lg n)$ time.

But destroying the input is kind of a drag—it might cause problems elsewhere in our code. Can we maintain this time and space cost without destroying the input?

Let's take a step back. **How can we break this problem down into subproblems?**

If we're going to do $O(n \lg n)$ time, we'll probably be iteratively doubling something or iteratively cutting something in half. That's how we usually get a " $\lg n$ ". So what if we could cut the problem in half somehow?

Well, binary search works by cutting the problem in half after figuring out which half of our input vector holds the answer.

But in a binary search, the *reason* we can confidently say which half has the answer is because the vector is *sorted*. For this problem, when we cut our unsorted vector in half we can't really make any strong statements about which elements are in the left half and which are in the right half.

What if we could cut the problem in half a *different* way, other than cutting the *vector* in half?

With this problem, we're looking for a needle (a repeated number) in a haystack (vector). What if instead of cutting the haystack in half, we cut *the set of possibilities for the needle* in half?

The full range of possibilities for our needle is $1..n$. How could we test whether the actual needle is in the first half of that range ($1..\frac{n}{2}$) or the second half ($\frac{n}{2} + 1..n$)?

A quick note about how we're defining our ranges: when we take $\frac{n}{2}$ we're doing *integer division*, so we throw away the remainder. To see what's going on, we should look at what happens when n is even and when n is odd:

- If n is 6 (an even number), we have $\frac{n}{2} = 3$ and $\frac{n}{2} + 1 = 4$, so our ranges are $1..3$ and $4..6$.
- If n is 5 (an odd number), $\frac{n}{2} = 2$ (we throw out the remainder) and $\frac{n}{2} + 1 = 3$, so our ranges are $1..2$ and $3..5$.

So we can notice a few properties about our ranges:

1. They aren't necessarily the same size.
2. They don't overlap.
3. Taken *together*, they represent the original input vector's range of $1..n$. In math terminology, we could say their *union* is $1..n$.

So, how do we know if the needle is in $1.. \frac{n}{2}$ or $\frac{n}{2} + 1..n$?

Think about the original problem statement. We know that we have at least one repeat because there are $n + 1$ items and they are all in the range $1..n$, which contains only n distinct integers.

This notion of "we have more items than we have possibilities, so we must have at least one repeat" is pretty powerful. It's sometimes called the pigeonhole principle. Can we exploit the pigeonhole principle to see which half of our range contains a repeat?

Imagine that we separated the input vector into two subvectors—one containing the items in the range $1.. \frac{n}{2}$ and the other containing the items in the range $\frac{n}{2} + 1..n$.

Each subvector has *a number of elements* as well as *a number of possible distinct integers* (that is, the length of the range of possible integers it holds).

Given what we know about the number of elements vs the number of possible distinct integers in the *original input vector*, what can we say about the number of elements vs the number of distinct possible integers in *these subvectors*?

The sum of the subvectors' numbers of elements is $n + 1$ (the number of elements in the original input vector) and the sum of the subvectors' numbers of possible distinct integers is n (the number of possible distinct integers in the original input vector).

Since the sums of the subvectors' numbers of elements must be 1 greater than the sum of the subvectors' numbers of possible distinct integers, one of the subvectors must have at least one more element than it has possible distinct integers.

Not convinced? We can prove this by contradiction. Suppose neither vector had more elements than it had possible distinct integers. In other words, both vectors have *at most* the same number of items as they have distinct possibilities. The sum of their numbers of items would then be *at most* the total number of possibilities across each of them, which is n . This is a contradiction—we know that our total number of items from the original input vector is $n + 1$, which is greater than n .

Now that we know *one* of our subvectors has 1 or more items more than it has distinct possibilities, we know *that subvector* must have at least one duplicate, by the same pigeonhole argument that we use to know that the *original input vector* has at least one duplicate.

So once we know *which* subvector has the count higher than its number of distinct possibilities, we can use this same approach recursively, cutting *that* subvector into two halves, etc, until we have just 1 item left in our range.

Of course, we don't need to actually separate our vector into subvectors. All we care about is *how long* each subvector would be. So we can simply do one walk through the input vector, counting the number of items that *would be* in each subvector.

Can you formalize this in code?

Careful—if we do this recursively, we'll incur a space cost in the call stack! Do it iteratively instead.

Solution

Our approach is similar to a binary search, except we divide the *range of possible answers* in half at each step, rather than dividing the *vector* in half.

1. Find the number of integers in our input vector which lie within the range $1.. \frac{n}{2}$.
2. Compare that to the number of possible unique integers in the same range.
3. If the number of *actual* integers is *greater* than the number of *possible* integers, we know there's a duplicate in the range $1.. \frac{n}{2}$, so we iteratively use the same approach on that range.
4. If the number of actual integers is *not greater* than the number of possible integers, we know there must be duplicate in the range $\frac{n}{2} + 1..n$, so we iteratively use the same approach on that range.
5. At some point, our range will contain just 1 integer, which will be our answer.

```
unsigned int findRepeat(const vector<unsigned int>& numbers)
{
    unsigned int floor = 1;
    unsigned int ceiling = numbers.size() - 1;

    while (floor < ceiling) {

        // divide our range 1..n into an upper range and lower range
        // (such that they don't overlap)
        // lower range is floor..midpoint
        // upper range is midpoint+1..ceiling
        unsigned int midpoint = floor + ((ceiling - floor) / 2);
        unsigned int lowerRangeFloor    = floor;
        unsigned int lowerRangeCeiling = midpoint;
        unsigned int upperRangeFloor    = midpoint + 1;
        unsigned int upperRangeCeiling = ceiling;

        // count number of items in lower range
        unsigned int itemsInLowerRange = 0;
        for (unsigned int item : numbers) {
            // is it in the lower range?
            if (item >= lowerRangeFloor && item <= lowerRangeCeiling) {
                ++itemsInLowerRange;
            }
        }

        unsigned int distinctPossibleIntegersInLowerRange = lowerRangeCeiling - lowerRangeFloor + 1;

        if (itemsInLowerRange > distinctPossibleIntegersInLowerRange) {
            // there must be a duplicate in the lower range
            // so use the same approach iteratively on that range
            floor    = lowerRangeFloor;
            ceiling = lowerRangeCeiling;
        }
        else {
            // there must be a duplicate in the upper range
            // so use the same approach iteratively on that range
            floor    = upperRangeFloor;
            ceiling = upperRangeCeiling;
        }
    }
}
```

```
}

// floor and ceiling have converged
// we found a number that repeats!
return floor;
}
```

Complexity

$O(1)$ space and $O(n \lg n)$ time.

Tricky as this solution is, we can actually do even better, getting our runtime down to $O(n)$ while keeping our space cost at $O(1)$. The solution is NUTS; it's probably outside the scope of what most interviewers would expect. But for the curious...here it is (</question/find-duplicate-optimize-for-space-beast-mode/>)!

Bonus

This function always returns *one* duplicate, but there may be several duplicates. Write a function that returns *all* duplicates.

What We Learned

Our answer was a modified binary search. We got there by *reasoning about the expected runtime*:

1. We started with an $O(n^2)$ "brute force" solution and wondered if we could do better.
2. We knew to beat $O(n^2)$ we'd probably do $O(n)$ or $O(n \lg n)$, so we started thinking of ways we might get an $O(n \lg n)$ runtime.
3. $\lg(n)$ usually comes from iteratively cutting stuff in half, so we arrived at the final algorithm by exploring that idea.

Starting with a target runtime and working *backward* from there can be a powerful strategy for all kinds of coding interview questions.

Ready for more?

Check out our full course →

Want more coding interview help?

Check out **interviewcake.com** for more advice, guides, and practice questions.